

Satellite Images Classification using Pretrained Models using Attention Mechanism



```
import numpy as np
import pandas as pd

image_dir = '/kaggle/input/remote-sensing-satellite-images/Remote
Sensing Data.v2i.yolov8/train/images'
label_dir = '/kaggle/input/remote-sensing-satellite-images/Remote
Sensing Data.v2i.yolov8/train/labels'

image_files = sorted([f for f in os.listdir(image_dir) if
f.endswith('.jpg') or f.endswith('.png')])
label_files = sorted([f for f in os.listdir(label_dir) if
f.endswith('.txt')])

image_paths = []
class_labels = []
center_x = []
center_y = []
width = []
height = []

for img, lbl in zip(image_files, label_files):
    img_path = os.path.join(image_dir, img)
```

```

label_path = os.path.join(label_dir, lbl)

with open(label_path, 'r') as file:
    for line in file:

        parts = line.strip().split()
        if len(parts) == 5:

            image_paths.append(img_path)
            class_labels.append(int(parts[0]))
            center_x.append(float(parts[1]))
            center_y.append(float(parts[2]))
            width.append(float(parts[3]))
            height.append(float(parts[4]))

df = pd.DataFrame({
    'image_path': image_paths,
    'class_label': class_labels,
    'center_x': center_x,
    'center_y': center_y,
    'width': width,
    'height': height
})

df.head()

```

	image_path	class_label
center_x \		
0	/kaggle/input/remote-sensing-satellite-images/...	2
0.498437		
1	/kaggle/input/remote-sensing-satellite-images/...	10
0.510938		
2	/kaggle/input/remote-sensing-satellite-images/...	9
0.498437		
3	/kaggle/input/remote-sensing-satellite-images/...	4
0.498437		
4	/kaggle/input/remote-sensing-satellite-images/...	5
0.498437		

```


```

	center_y	width	height
0	0.498437	0.996094	0.996094
1	0.618750	0.970313	0.754687
2	0.498437	0.996094	0.996094
3	0.498437	0.996094	0.996094
4	0.498437	0.996094	0.996094

```

df.tail()

```

	image_path	class_label
center_x \		

```

743 /kaggle/input/remote-sensing-satellite-images/... 7
0.498437
744 /kaggle/input/remote-sensing-satellite-images/... 3
0.498437
745 /kaggle/input/remote-sensing-satellite-images/... 2
0.498437
746 /kaggle/input/remote-sensing-satellite-images/... 3
0.182031
747 /kaggle/input/remote-sensing-satellite-images/... 11
0.662500

```

```

      center_y    width    height
743  0.498437  0.996094  0.996094
744  0.498437  0.996094  0.996094
745  0.498437  0.996094  0.996094
746  0.498437  0.364063  0.996094
747  0.498437  0.667969  0.996094

```

```
df.shape
```

```
(748, 6)
```

```
df.columns
```

```
Index(['image_path', 'class_label', 'center_x', 'center_y', 'width',
      'height'], dtype='object')
```

```
df.duplicated().sum()
```

```
0
```

```
df.isnull().sum()
```

```

image_path    0
class_label   0
center_x      0
center_y      0
width         0
height        0
dtype: int64

```

```
df = df[['image_path', 'class_label']]
```

```
df
```

```

      image_path    class_label
0  /kaggle/input/remote-sensing-satellite-images/...      2
1  /kaggle/input/remote-sensing-satellite-images/...     10
2  /kaggle/input/remote-sensing-satellite-images/...      9
3  /kaggle/input/remote-sensing-satellite-images/...      4
4  /kaggle/input/remote-sensing-satellite-images/...      5
..          ...          ...

```

```
743 /kaggle/input/remote-sensing-satellite-images/... 7
744 /kaggle/input/remote-sensing-satellite-images/... 3
745 /kaggle/input/remote-sensing-satellite-images/... 2
746 /kaggle/input/remote-sensing-satellite-images/... 3
747 /kaggle/input/remote-sensing-satellite-images/... 11
```

```
[748 rows x 2 columns]
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

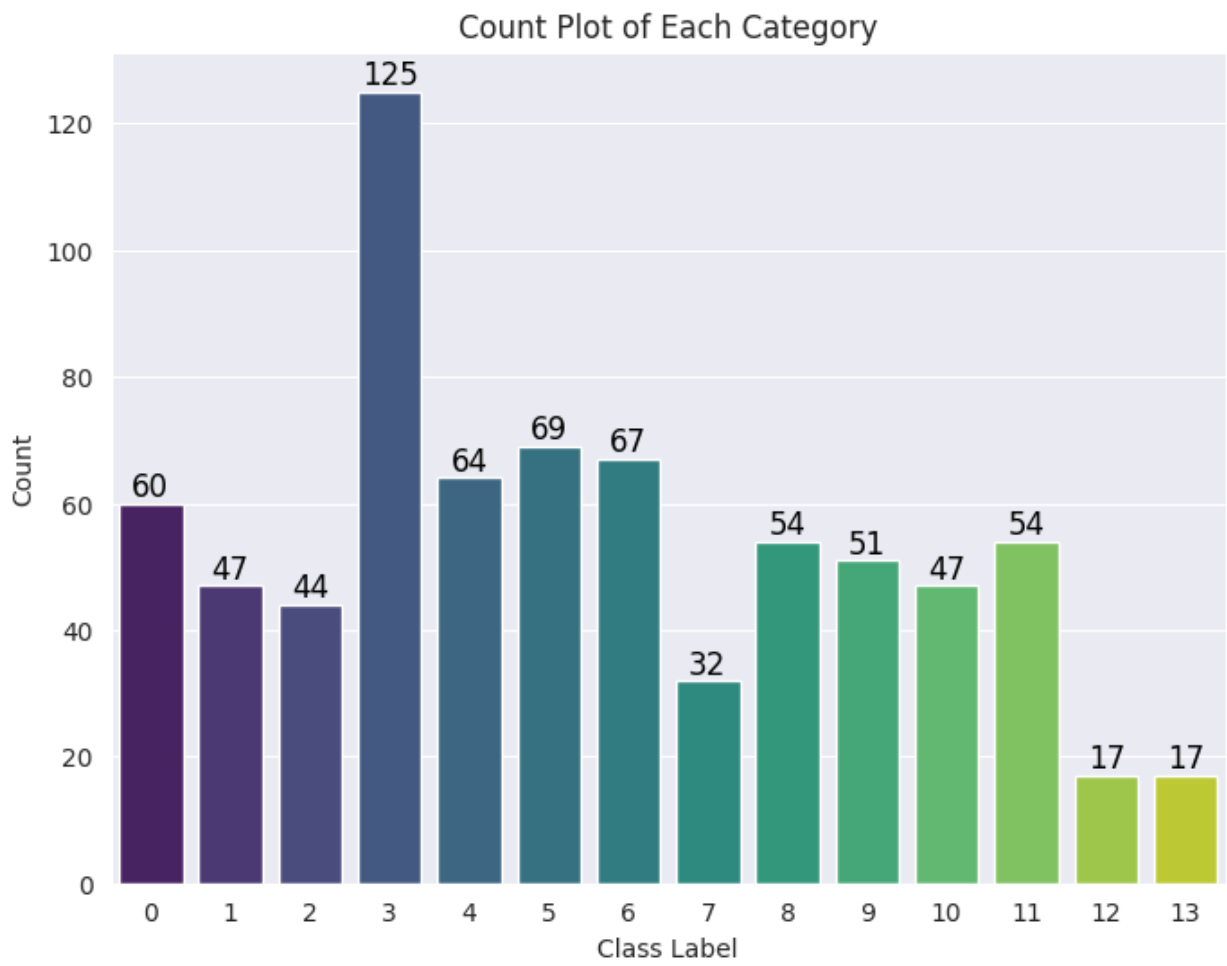
```
plt.figure(figsize=(8, 6))
ax = sns.countplot(data=df, x='class_label', palette="viridis")
plt.title('Count Plot of Each Category')
plt.xlabel('Class Label')
plt.ylabel('Count')

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}',
                (p.get_x() + p.get_width() / 2, p.get_height()),
                ha='center', va='bottom', fontsize=12, color='black')
```


```
plt.show()
```

```
plt.figure(figsize=(8, 8))
class_counts = df['class_label'].value_counts()
plt.pie(class_counts, labels=class_counts.index, autopct='%1.1f%%',
        colors=sns.color_palette("viridis"), startangle=140)
```

```
plt.title('Pie Chart of Each Category')
plt.show()
```



Pie Chart of Each Category



```
import random
from PIL import Image

images_per_category = 5

class_labels = df['class_label'].unique()

fig, axes = plt.subplots(len(class_labels), images_per_category,
figsize=(15, 3 * len(class_labels)))

for i, class_label in enumerate(class_labels):
```

```
class_images = df[df['class_label'] == class_label]
['image_path'].tolist()


sampled_images = random.sample(class_images,
min(images_per_category, len(class_images)))

for j, image_path in enumerate(sampled_images):
    img = Image.open(image_path)

    axes[i, j].imshow(img)
    axes[i, j].axis('off')

    if j == 0:
        axes[i, j].set_ylabel(f"Class {class_label}", rotation=0,
labelpad=50, fontsize=12, ha='right', va='center')

plt.tight_layout()
plt.show()
```


```

from imblearn.over_sampling import RandomOverSampler

ros = RandomOverSampler(random_state=42)
X_resampled, y_resampled = ros.fit_resample(df[['image_path']],
df['class_label'])

df_resampled = pd.DataFrame(X_resampled, columns=['image_path'])
df_resampled['class_label'] = y_resampled

print("\nClass distribution after oversampling:")
print(df_resampled['class_label'].value_counts())

```

Class distribution after oversampling:

class_label

2 125

10 125

9 125

4 125

5 125

8 125

1 125

13 125

0 125

3 125

11 125

6 125

7 125

12 125

Name: count, dtype: int64

df_resampled

	image_path	class_label
0	/kaggle/input/remote-sensing-satellite-images/...	2
1	/kaggle/input/remote-sensing-satellite-images/...	10
2	/kaggle/input/remote-sensing-satellite-images/...	9
3	/kaggle/input/remote-sensing-satellite-images/...	4
4	/kaggle/input/remote-sensing-satellite-images/...	5
...
1745	/kaggle/input/remote-sensing-satellite-images/...	13
1746	/kaggle/input/remote-sensing-satellite-images/...	13
1747	/kaggle/input/remote-sensing-satellite-images/...	13
1748	/kaggle/input/remote-sensing-satellite-images/...	13
1749	/kaggle/input/remote-sensing-satellite-images/...	13

[1750 rows x 2 columns]

```

import time
import shutil
import pathlib

```

```

import itertools
from PIL import Image

import cv2
import seaborn as sns
sns.set_style('darkgrid')
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam, Adamax
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten,
Dense, Activation, Dropout, BatchNormalization
from tensorflow.keras import regularizers

import warnings
warnings.filterwarnings("ignore")

print ('check')

check

df_resampled['class_label'] = df_resampled['class_label'].astype(str)
train_df_new, temp_df_new = train_test_split(
    df_resampled,
    train_size=0.8,
    shuffle=True,
    random_state=42,
    stratify=df_resampled['class_label']
)

valid_df_new, test_df_new = train_test_split(
    temp_df_new,
    test_size=0.5,
    shuffle=True,
    random_state=42,
    stratify=temp_df_new['class_label']
)

from tensorflow.keras.preprocessing.image import ImageDataGenerator

batch_size = 16
img_size = (224, 224)
channels = 3
img_shape = (img_size[0], img_size[1], channels)

```

```
tr_gen = ImageDataGenerator(rescale=1./255)
ts_gen = ImageDataGenerator(rescale=1./255)
```

```
train_gen_new = tr_gen.flow_from_dataframe(
    train_df_new,
    x_col='image_path',
    y_col='class_label',
    target_size=img_size,
    class_mode='sparse',
    color_mode='rgb',
    shuffle=True,
    batch_size=batch_size
)
```

```
valid_gen_new = ts_gen.flow_from_dataframe(
    valid_df_new,
    x_col='image_path',
    y_col='class_label',
    target_size=img_size,
    class_mode='sparse',
    color_mode='rgb',
    shuffle=True,
    batch_size=batch_size
)
```

```
test_gen_new = ts_gen.flow_from_dataframe(
    test_df_new,
    x_col='image_path',
    y_col='class_label',
    target_size=img_size,
    class_mode='sparse',
    color_mode='rgb',
    shuffle=False,
    batch_size=batch_size
)
```

Found 1400 validated image filenames belonging to 14 classes.

Found 175 validated image filenames belonging to 14 classes.

Found 175 validated image filenames belonging to 14 classes.

```
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint

physical_devices = tf.config.list_physical_devices('GPU')
if physical_devices:
    print("Using GPU")
else:
    print("Using CPU")
```

Using GPU

```
early_stopping = EarlyStopping(monitor='val_loss', patience=5,
restore_best_weights=True)

from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense,
Dropout, BatchNormalization, GaussianNoise, Input, MultiHeadAttention,
Reshape
from tensorflow.keras.optimizers import Adam
import tensorflow as tf

def create_vgg16_model(input_shape):

    inputs = Input(shape=input_shape)

    base_model = VGG16(weights='imagenet', input_tensor=inputs,
include_top=False)

    for layer in base_model.layers:
        layer.trainable = False

    x = base_model.output

    height, width, channels = 7, 7, 512
    x = Reshape((height * width, channels))(x)

    attention_output = MultiHeadAttention(num_heads=8,
key_dim=channels)(x, x)

    attention_output = Reshape((height, width, channels))
(attention_output)

    x = GaussianNoise(0.25)(attention_output)
    x = GlobalAveragePooling2D()(x)

    x = Dense(512, activation='relu')(x)
    x = BatchNormalization()(x)
    x = GaussianNoise(0.25)(x)
    x = Dropout(0.25)(x)

    outputs = Dense(14, activation='softmax')(x)

    model = Model(inputs=inputs, outputs=outputs)

    return model

input_shape = (224, 224, 3)
cnn_model = create_vgg16_model(input_shape)
```

```

cnn_model.compile(optimizer=Adam(learning_rate=0.0001),
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])


history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=5,
    callbacks=[early_stopping],
    verbose=1
)


Epoch 1/5
88/88 ━━━━━━━━━━━ 23s 167ms/step - accuracy: 0.3485 - loss:
2.0396 - val_accuracy: 0.3886 - val_loss: 2.2730
Epoch 2/5
88/88 ━━━━━━━━━━━ 8s 90ms/step - accuracy: 0.6384 - loss:
1.2135 - val_accuracy: 0.4914 - val_loss: 1.9934
Epoch 3/5
88/88 ━━━━━━━━━━━ 8s 91ms/step - accuracy: 0.6936 - loss:
0.9943 - val_accuracy: 0.6800 - val_loss: 1.5046
Epoch 4/5
88/88 ━━━━━━━━━━━ 8s 91ms/step - accuracy: 0.7599 - loss:
0.7450 - val_accuracy: 0.6914 - val_loss: 1.1762
Epoch 5/5
88/88 ━━━━━━━━━━━ 8s 90ms/step - accuracy: 0.8017 - loss:
0.6267 - val_accuracy: 0.6629 - val_loss: 1.0204

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

```



```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_classes = np.argmax(predictions, axis=1)
```

11/11 ————— 2s 125ms/step

```
report = classification_report(test_labels, predicted_classes,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
0	1.00	0.25	0.40	12
1	0.79	0.85	0.81	13
10	0.60	1.00	0.75	12
11	0.77	0.77	0.77	13
12	0.85	0.85	0.85	13
13	0.50	0.58	0.54	12
2	1.00	0.85	0.92	13
3	1.00	0.38	0.56	13
4	0.35	0.92	0.51	12
5	0.75	0.50	0.60	12
6	0.89	0.62	0.73	13


7	0.92	1.00	0.96	12
8	0.64	0.58	0.61	12
9	0.80	0.62	0.70	13
accuracy			0.70	175
macro avg	0.78	0.70	0.69	175
weighted avg	0.78	0.70	0.70	175

```

conf_matrix = confusion_matrix(test_labels, predicted_classes)

plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=list(test_gen_new.class_indices.keys()),
            yticklabels=list(test_gen_new.class_indices.keys()))
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

```



```

from tensorflow.keras.applications import VGG19
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense,
Dropout, BatchNormalization, GaussianNoise, Input, MultiHeadAttention,
Reshape
from tensorflow.keras.optimizers import Adam
import tensorflow as tf

def create_vgg19_model(input_shape):
    inputs = Input(shape=input_shape)

    base_model = VGG19(weights='imagenet', input_tensor=inputs,
include_top=False)

    for layer in base_model.layers:

```

```

        layer.trainable = False

    x = base_model.output

    height, width, channels = 7, 7, 512
    x = Reshape((height * width, channels))(x)

    attention_output = MultiHeadAttention(num_heads=8,
key_dim=channels)(x, x)

    attention_output = Reshape((height, width, channels))
(attention_output)

    x = GaussianNoise(0.25)(attention_output)
    x = GlobalAveragePooling2D()(x)

    x = Dense(512, activation='relu')(x)
    x = BatchNormalization()(x)
    x = GaussianNoise(0.25)(x)
    x = Dropout(0.25)(x)

    outputs = Dense(14, activation='softmax')(x)

    model = Model(inputs=inputs, outputs=outputs)

    return model

input_shape = (224, 224, 3)
cnn_model = create_vgg19_model(input_shape)

cnn_model.compile(optimizer=Adam(learning_rate=0.0001),
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=5,
    callbacks=[early_stopping],
    verbose=1
)

Epoch 1/5
88/88 _____ 24s 179ms/step - accuracy: 0.3077 - loss:
2.1396 - val_accuracy: 0.2800 - val_loss: 2.3190
Epoch 2/5
88/88 _____ 9s 102ms/step - accuracy: 0.5888 - loss:
1.2940 - val_accuracy: 0.2743 - val_loss: 2.1177
Epoch 3/5
88/88 _____ 10s 105ms/step - accuracy: 0.6390 - loss:
1.1083 - val_accuracy: 0.4057 - val_loss: 1.7282

```



Epoch 4/5
88/88 10s 107ms/step - accuracy: 0.7429 - loss: 0.8452 - val_accuracy: 0.5886 - val_loss: 1.3326

Epoch 5/5
88/88 10s 107ms/step - accuracy: 0.7924 - loss: 0.6552 - val_accuracy: 0.5829 - val_loss: 1.3261

```
plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])  
plt.title('Model accuracy')  
plt.ylabel('Accuracy')  
plt.xlabel('Epoch')  
plt.legend(['Train', 'Validation'], loc='upper left')  
plt.show()
```

```
plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])  
plt.title('Model loss')  
plt.ylabel('Loss')  
plt.xlabel('Epoch')  
plt.legend(['Train', 'Validation'], loc='upper left')  
plt.show()
```





```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_classes = np.argmax(predictions, axis=1)
```

11/11 ————— 2s 142ms/step

```
report = classification_report(test_labels, predicted_classes,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
0	1.00	0.42	0.59	12
1	0.40	0.15	0.22	13
10	0.45	0.42	0.43	12
11	0.21	0.92	0.34	13
12	0.67	0.15	0.25	13
13	0.00	0.00	0.00	12
2	1.00	0.08	0.14	13
3	0.00	0.00	0.00	13
4	0.00	0.00	0.00	12
5	0.00	0.00	0.00	12
6	0.00	0.00	0.00	13


7	1.00	0.33	0.50	12
8	0.13	0.83	0.22	12
9	0.90	0.69	0.78	13
accuracy			0.29	175
macro avg	0.41	0.29	0.25	175
weighted avg	0.41	0.29	0.25	175

```

conf_matrix = confusion_matrix(test_labels, predicted_classes)

plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=list(test_gen_new.class_indices.keys()),
            yticklabels=list(test_gen_new.class_indices.keys()))
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

```



```

from tensorflow.keras.applications import InceptionV3
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense,
Dropout, BatchNormalization, GaussianNoise, Input, MultiHeadAttention,
Reshape
from tensorflow.keras.optimizers import Adam
import tensorflow as tf

def create_inceptionv3_model(input_shape):
    inputs = Input(shape=input_shape)

    base_model = InceptionV3(weights='imagenet', input_tensor=inputs,
include_top=False)

    for layer in base_model.layers:

```

```

        layer.trainable = False

    x = base_model.output

    height, width, channels = 5, 5, 2048
    x = Reshape((height * width, channels))(x)

    attention_output = MultiHeadAttention(num_heads=8,
key_dim=channels)(x, x)

    attention_output = Reshape((height, width, channels))
(attention_output)

    x = GaussianNoise(0.25)(attention_output)
    x = GlobalAveragePooling2D()(x)

    x = Dense(512, activation='relu')(x)
    x = BatchNormalization()(x)
    x = GaussianNoise(0.25)(x)
    x = Dropout(0.25)(x)

    outputs = Dense(14, activation='softmax')(x)

    model = Model(inputs=inputs, outputs=outputs)

    return model

input_shape = (224, 224, 3)
cnn_model = create_inceptionv3_model(input_shape)

cnn_model.compile(optimizer=Adam(learning_rate=0.0001),
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])

history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=5,
    callbacks=[early_stopping],
    verbose=1
)

Epoch 1/5
88/88 _____ 48s 340ms/step - accuracy: 0.4807 - loss:
1.9312 - val_accuracy: 0.5657 - val_loss: 1.7963
Epoch 2/5
88/88 _____ 12s 129ms/step - accuracy: 0.8081 - loss:
0.6570 - val_accuracy: 0.5886 - val_loss: 1.4472
Epoch 3/5
88/88 _____ 12s 131ms/step - accuracy: 0.8700 - loss:
0.4519 - val_accuracy: 0.6629 - val_loss: 1.2109

```



Epoch 4/5
88/88 12s 131ms/step - accuracy: 0.8957 - loss: 0.3675 - val_accuracy: 0.5314 - val_loss: 2.7913

Epoch 5/5
88/88 12s 129ms/step - accuracy: 0.9102 - loss: 0.3699 - val_accuracy: 0.4286 - val_loss: 3.7984

```
plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])  
plt.title('Model accuracy')  
plt.ylabel('Accuracy')  
plt.xlabel('Epoch')  
plt.legend(['Train', 'Validation'], loc='upper left')  
plt.show()
```

```
plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])  
plt.title('Model loss')  
plt.ylabel('Loss')  
plt.xlabel('Epoch')  
plt.legend(['Train', 'Validation'], loc='upper left')  
plt.show()
```





```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_classes = np.argmax(predictions, axis=1)
```

11/11 ————— 11s 570ms/step

```
report = classification_report(test_labels, predicted_classes,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
0	1.00	0.08	0.15	12
1	0.62	0.77	0.69	13
10	0.79	0.92	0.85	12
11	0.76	1.00	0.87	13
12	0.88	0.54	0.67	13
13	0.89	0.67	0.76	12
2	0.33	1.00	0.49	13
3	0.34	0.85	0.49	13
4	1.00	0.08	0.15	12
5	0.50	0.08	0.14	12
6	0.86	0.46	0.60	13


7	0.00	0.00	0.00	12
8	0.50	0.42	0.45	12
9	0.67	0.92	0.77	13
accuracy			0.57	175
macro avg	0.65	0.56	0.51	175
weighted avg	0.65	0.57	0.51	175

```

conf_matrix = confusion_matrix(test_labels, predicted_classes)

plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=list(test_gen_new.class_indices.keys()),
            yticklabels=list(test_gen_new.class_indices.keys()))
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

```



```

from tensorflow.keras.applications import Xception
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense,
Dropout, BatchNormalization, GaussianNoise, Input, MultiHeadAttention,
Reshape
from tensorflow.keras.optimizers import Adam
import tensorflow as tf

def create_xception_model(input_shape):
    inputs = Input(shape=input_shape)

    base_model = Xception(weights='imagenet', input_tensor=inputs,
include_top=False)

    for layer in base_model.layers:

```

```

        layer.trainable = False

    x = base_model.output

    height, width, channels = 7, 7, 2048
    x = Reshape((height * width, channels))(x)

    attention_output = MultiHeadAttention(num_heads=8,
key_dim=channels)(x, x)

    attention_output = Reshape((height, width, channels))
(attention_output)

    x = GaussianNoise(0.25)(attention_output)
    x = GlobalAveragePooling2D()(x)
    x = Dense(512, activation='relu')(x)
    x = BatchNormalization()(x)
    x = GaussianNoise(0.25)(x)
    x = Dropout(0.25)(x)
    outputs = Dense(14, activation='softmax')(x)

    model = Model(inputs=inputs, outputs=outputs)

    return model

input_shape = (224, 224, 3)
cnn_model = create_xception_model(input_shape)

cnn_model.compile(optimizer=Adam(learning_rate=0.0001),
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])

history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=5,
    callbacks=[early_stopping],
    verbose=1
)

Epoch 1/5
88/88 _____ 43s 345ms/step - accuracy: 0.5023 - loss:
1.7452 - val_accuracy: 0.5200 - val_loss: 1.5167
Epoch 2/5
88/88 _____ 21s 238ms/step - accuracy: 0.8033 - loss:
0.6875 - val_accuracy: 0.7600 - val_loss: 0.8965
Epoch 3/5
88/88 _____ 21s 232ms/step - accuracy: 0.8390 - loss:
0.5244 - val_accuracy: 0.8000 - val_loss: 0.7940
Epoch 4/5

```





```
88/88 _____ 20s 226ms/step - accuracy: 0.8884 - loss: 0.3722 - val_accuracy: 0.7714 - val_loss: 0.7607  
Epoch 5/5
```

```
88/88 _____ 20s 218ms/step - accuracy: 0.8929 - loss: 0.3921 - val_accuracy: 0.7657 - val_loss: 0.7880
```

```
plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])  
plt.title('Model accuracy')  
plt.ylabel('Accuracy')  
plt.xlabel('Epoch')  
plt.legend(['Train', 'Validation'], loc='upper left')  
plt.show()
```

```
plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])  
plt.title('Model loss')  
plt.ylabel('Loss')  
plt.xlabel('Epoch')  
plt.legend(['Train', 'Validation'], loc='upper left')  
plt.show()
```





```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_classes = np.argmax(predictions, axis=1)
```

11/11 ————— 6s 335ms/step

```
report = classification_report(test_labels, predicted_classes,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
0	0.75	0.75	0.75	12
1	1.00	0.85	0.92	13
10	0.79	0.92	0.85	12
11	0.92	0.92	0.92	13
12	1.00	1.00	1.00	13
13	0.80	0.67	0.73	12
2	1.00	1.00	1.00	13
3	0.75	0.69	0.72	13
4	1.00	0.67	0.80	12
5	0.55	0.50	0.52	12
6	0.46	0.92	0.62	13


7	0.86	0.50	0.63	12
8	0.75	0.50	0.60	12
9	0.71	0.92	0.80	13
accuracy			0.78	175
macro avg	0.81	0.77	0.78	175
weighted avg	0.81	0.78	0.78	175

```

conf_matrix = confusion_matrix(test_labels, predicted_classes)

plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=list(test_gen_new.class_indices.keys()),
            yticklabels=list(test_gen_new.class_indices.keys()))
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

```



```
import matplotlib.pyplot as plt
import numpy as np

labels = [str(i) for i in range(10)]

vgg16_precision = [1.00, 0.79, 0.60, 0.77, 0.85, 0.50, 1.00, 1.00,
0.35, 0.75]
vgg16_recall = [0.25, 0.85, 1.00, 0.77, 0.85, 0.58, 0.85, 0.38, 0.92,
0.50]
vgg16_f1 = [0.40, 0.81, 0.75, 0.77, 0.85, 0.54, 0.92, 0.56, 0.51,
0.60]

vgg19_precision = [1.00, 0.40, 0.45, 0.21, 0.67, 0.00, 1.00, 0.00,
0.00, 0.90]
vgg19_recall = [0.42, 0.15, 0.42, 0.92, 0.15, 0.00, 0.08, 0.00, 0.00,
0.69]
```

```

vgg19_f1 = [0.59, 0.22, 0.43, 0.34, 0.25, 0.00, 0.14, 0.00, 0.00,
0.78]

inceptionv3_precision = [1.00, 0.62, 0.79, 0.76, 0.88, 0.89, 0.33,
0.34, 1.00, 0.50]
inceptionv3_recall = [0.08, 0.77, 0.92, 1.00, 0.54, 0.67, 1.00, 0.85,
0.08, 0.42]
inceptionv3_f1 = [0.15, 0.69, 0.85, 0.87, 0.67, 0.76, 0.49, 0.49,
0.15, 0.77]

xception_precision = [0.75, 1.00, 0.79, 0.92, 1.00, 0.80, 1.00, 0.75,
1.00, 0.71]
xception_recall = [0.75, 0.85, 0.92, 0.92, 1.00, 0.67, 1.00, 0.69,
0.67, 0.92]
xception_f1 = [0.75, 0.92, 0.85, 0.92, 1.00, 0.73, 1.00, 0.72, 0.80,
0.80]

fig, ax = plt.subplots(1, 3, figsize=(18, 6))

models = ['VGG16', 'VGG19', 'InceptionV3', 'Xception']
precision_data = [vgg16_precision, vgg19_precision,
inceptionv3_precision, xception_precision]
recall_data = [vgg16_recall, vgg19_recall, inceptionv3_recall,
xception_recall]
f1_data = [vgg16_f1, vgg19_f1, inceptionv3_f1, xception_f1]

x = np.arange(len(labels))
width = 0.2

for i, model in enumerate(models):
    ax[0].bar(x + i * width - 1.5 * width, precision_data[i], width,
label=model)


for i, model in enumerate(models):
    ax[1].bar(x + i * width - 1.5 * width, recall_data[i], width,
label=model)

for i, model in enumerate(models):
    ax[2].bar(x + i * width - 1.5 * width, f1_data[i], width,
label=model)

for i in range(3):
    ax[i].set_xlabel('Classes')
    ax[i].set_ylabel('Scores')
    ax[i].set_title(f'{models[i]} Results')
    ax[i].set_xticks(x)
    ax[i].set_xticklabels(labels)
    ax[i].legend()

```

```
plt.tight_layout()
plt.show()
```



Thanks !!!